Recursive Multi-Agent Mixture-of-Experts (R-MoE) for Autonomous Clinical Diagnostics

The introduction of large-scale vision-language models (VLMs) into the clinical environment has marked a fundamental transition in how medical imaging is interpreted, summarized, and integrated into broader patient care pathways. Historically, the application of artificial intelligence in radiology was limited to narrow tasks such as bone age determination or basic nodule detection using convolutional neural networks. However, the emergence of the Recursive Multi-Agent Mixture-of-Experts (R-MoE) framework represents a more sophisticated architectural shift, moving away from monolithic, single-pass inference models toward modular, agentic systems that mimic the dual-process cognitive workflows of human radiologists.[1] This report provides an exhaustive analysis of the R-MoE architecture, detailing its reliance on the Qwen2-VL-72B perception engine, the DeepSeek-R1 reasoning layer, and the Llama-3-Medius clinical reporting specialist, while exploring the novel #wanna# recursive protocol as a primary mechanism for diagnostic safety.[4]

Architectural Foundations of the Multi-Modal Perception Engine

The first phase of the R-MoE framework is the Multi-Modal Perception Engine (MPE), which utilizes the Qwen2-VL-72B model as its primary expert. The MPE is tasked with global feature extraction, lesion segmentation, spatial relationship mapping, and anatomical landmark identification.[4] One of the most critical challenges in medical imaging is the high variability in image resolution and aspect ratio, particularly when comparing standard chest X-rays to high-resolution CT slices or specialized MRI sequences.[4] Qwen2-VL-72B addresses this through the Naive Dynamic Resolution mechanism, which allows the model to dynamically process images of varying resolutions into a flexible number of visual tokens.[4] This ensures that subtle pathological features, which might be lost in fixed-resolution architectures, are preserved for high-dimensional latent representation.[4]

The perceptual accuracy of the MPE is further enhanced by the integration of Multimodal Rotary Position Embedding (M-RoPE). This mechanism facilitates the fusion of positional information across text and visual data, which is essential for multi-view alignment in radiological series.[4] For instance, when analyzing a standard three-view elbow series for pediatric fractures, the MPE must maintain a consistent spatial context across the lateral, oblique, and anteroposterior views.[4] The MPE also performs artifact filtering and denoising, which is critical in clinical settings where patient movement or technical limitations can

introduce non-pathological signals that generalist models often misinterpret as disease.[4]

VLM Perceptual Performance	PathMMU Avg (%)	CT (%)	MRI (%)	X-Ray (%)
Qwen2-VL-72B	64.0	75.1	68.6	68.2
InternVL2.5-78B	62.5	71.2	67.2	71.8
Qwen2-VL-7B	55.4	62.6	57.2	56.9
LLaVA-v1.5-13B	39.3	45.9	37.3	39.7

The data indicates that the 72B parameter scale of the MPE provides a significant advantage in identifying anatomical landmarks and segmenting lesions across diverse modalities.[10] However, research also reveals that a significant gap—approximately 65% in localization tasks—remains when comparing generalist VLMs to human experts on open-world clinical benchmarks like NOVA.[11] This gap underscores the necessity of transitioning from pure perception to an agentic reasoning layer that can audit visual findings against established clinical logic.[5]

Cognitive Modeling in the Agentic Reasoning and Logic Layer

Phase 2 of the framework, the Agentic Reasoning & Logic Layer (ARLL), is powered by DeepSeek-R1. This component is designed to emulate System 2 reasoning—the slow, analytical, and deliberate cognitive process required for complex clinical problem-solving.[5] While the MPE handles the rapid, intuitive "System 1" perception of patterns, the ARLL performs Chain-of-Thought (CoT) auditing to verify if the visual anomalies detected in Phase 1 match the patient's clinical history and laboratory findings.[5]

The ARLL achieves this through an explicit inference chain that decomposes the diagnostic problem into sequential decision points.[12] On the MedQA dataset, which utilizes clinical vignettes in the USMLE format, DeepSeek-R1 demonstrated a 93% diagnostic accuracy, exhibiting patterns of medical reasoning comparable to domain experts.[5] This includes the ability to generate a differential diagnosis (DDx) by weighing the probabilistic significance of competing hypotheses.[13]

Error Patterns and Bias Mitigation

Despite high accuracy, the ARLL is susceptible to specific cognitive biases that are well-documented in human clinical practice. The most prevalent of these is anchoring bias, where the model focuses on a classic imaging pattern while ignoring conflicting temporal data.[5] For example, in cases of suspected neonatal bowel obstruction, the model might anchor on a "double bubble" sign indicating duodenal atresia while failing to account for a delayed clinical presentation that points toward a more complex diagnosis like annular pancreas.[12]

Reasoning Error Classification	Frequency in MedQA Analysis (%)	Clinical Impact
Anchoring Bias	14.3	Risk of delayed surgical intervention
Difficulty with Conflicting Data	21.4	Inconsistent treatment recommendations
Limited Alternative Consideration	28.6	Missed secondary pathologies
Overthinking / Rationalization	35.7	Increased probability of hallucination

Research suggests a strong correlation between the length of the reasoning chain and the probability of error.[13] Extended explanations often signal latent uncertainty or an attempt to rationalize a conclusion that is not fully supported by the primary data.[13] To address this, the R-MoE framework integrates uncertainty quantification (UQ) metrics that monitor the entropy of the reasoning process in real-time.[15]

The Recursive #wanna# Protocol and Stochastic Gating

The core innovation of the R-MoE system is the implementation of a recursive re-evaluation loop triggered by the #wanna# protocol.[7] When the ARLL detects that the confidence score (S_c) for a diagnosis is below a predetermined safety threshold ($S_c < 0.90$), it generates the #wanna# token [User Query]. This token acts as a programmatic signal to the orchestration layer to suspend the report generation process and re-initiate perception with specific

constraints.[18]

Mathematical Mechanism of Uncertainty Triggers

The framework utilizes predictive entropy to estimate the reliability of each diagnostic step. The entropy is calculated based on the probability distribution over the top visual and textual tokens generated by the models.[16] In high-risk clinical scenarios, such as the detection of subtle radiographic fractures or intracranial hemorrhages, the system adopts an "Abstention Logic" approach.[15]

The stochastic gating mechanism determines whether to route the latent visual embeddings to the synthesis agent or back to the perception engine for further investigation.[2] This recursive process allows the system to:

1. **Request a Re-scan:** Command the MPE to re-evaluate specific regions of interest (ROI) with higher token density to detect subtle features.[4]
2. **Request Additional Modalities:** Identify that a primary X-ray is insufficient for the clinical question and suggest escalation to CT or MRI.[7]
3. **Perform Temporal Analysis:** Trigger a search for older imaging records to perform a comparative analysis of lesion progression over time.[15]

This loop continues until the confidence score exceeds the $S_c \geq 0.90$ threshold or until it is determined that the available data is insufficient, at which point the case is escalated to a human radiologist.[15]

Synthesis Agent and Structured Clinical Reporting

In Phase 3, the Clinical Synthesis & Reporting (CSR) agent, powered by Llama-3-Medius, takes the validated diagnostic logic stream and converts it into a professional, structured report.[6] This expert is fine-tuned to use standardized medical terminology, including SNOMED CT and ICD-11 classifications, ensuring that the AI's output is interoperable with hospital electronic health record (EHR) systems.[6]

Risk Stratification and Automated Recommendations

The CSR agent is trained to apply specific clinical risk stratification frameworks, such as the BI-RADS for breast imaging or TIRADS for thyroid nodules.[6] By integrating medical knowledge graphs (KGs), the Llama-3-Medius model can ground its findings in evidence-based guidelines, reducing the risk of purely generative hallucinations.[6]

Reporting Accuracy Metrics	KG-LLM (Llama-3-Med)	Llama-2 Baseline	GPT-4V (Zero-Shot)

F1 Score (General)	0.914	0.867	0.824
Top-1 Drug Accuracy	0.782	0.716	0.645
Unsafe Suggestions Reduction	50%	-	12%

The architecture of the Synthesis Agent includes a dual-layer validation mechanism.[6] The first layer is a semantic parser that extracts entities from the reasoning logic, while the second layer is a deterministic rule-checker that evaluates the final report against clinical safety protocols, such as drug-dose-duration rules for pediatric patients.[27] This ensures that the automated treatment recommendations provided in the "Impression" section of the report are both medically sound and legally compliant.[27]

Human-in-the-Loop Validation and Clinical Integration

The R-MoE framework concludes with a Human-in-the-Loop (HITL) review and approval stage.[26] This process recognizes that while AI can significantly speed up repetitive interpretative tasks, the final accountability remains with the human practitioner.[31] The HITL interface presents the radiologist with the final autonomous report alongside the "thinking trace" generated by DeepSeek-R1 and any recursive re-evaluation flags.[7]

Optimization of Radiologist Workflow

By utilizing confidence threshold filtering, the system allows radiologists to focus their cognitive resources on high-uncertainty cases.[26] For routine screenings where the R-MoE achieves a very high confidence score (e.g., $S_c > 0.98$), the HITL review becomes a rapid verification task, significantly increasing the throughput of the radiology department.[26] Studies have shown that this collaborative approach can improve healthcare diagnostics from an 92% accuracy rate for AI alone to an 99.5% accuracy rate when integrated with human expertise.[26]

Furthermore, the framework supports clinical education and patient communication. Llama-3-Medius can generate simplified "plain language" versions of the radiology reports for patients, translating complex terms into understandable summaries of their health status.[28] It also provides Socratic tutoring for medical students by highlighting the reasoning steps taken to reach a specific diagnosis.[30]

Technical and Operational Challenges in Recursive

MoE Systems

The deployment of a multi-agent recursive system like R-MoE is not without significant technical challenges. The primary concern is the computational overhead associated with maintaining three large-scale models (72B, 671B, and 70B parameters) and the latency introduced by recursive loops.[8]

Inference Efficiency and Scaling

To maintain real-time performance, the R-MoE system must utilize advanced MoE operators such as FlashMoE, which fuses expert computation and inter-GPU communication to reduce launch overheads and idle gaps.[34] Additionally, the use of parameter-efficient fine-tuning (PEFT) and quantization techniques for the gating layers allows the system to operate on internal hospital hardware, addressing data privacy and residency requirements.[30]

The framework also faces challenges in "hidden state" transfer between agents.[36] When the perception engine sends its output to the reasoning layer, there is a risk that critical spatial or intensity information is suppressed during the tokenization process.[36] To mitigate this, researchers are exploring the use of Multi-Modal Contextual Vectors (MCV) that preserve the last hidden states of the perception model across multiple layers, serving as latent instructions for the reasoning agent.[37]

Operational Challenge	Technical Mitigation Strategy	Impact on Deployment
Inference Latency	FlashMoE & Persistent GPU Kernels	6x lower latency, 9x higher utilization [34]
Data Privacy	On-Premise Open-Source Weights	Compliance with EU AI Act and GDPR [30]
Model Drift	Continuous HITL Feedback Loops	Sustained accuracy over clinical cycles [30]
Information Loss	Multi-Modal Contextual Vectors	Enhanced spatial grounding in reasoning [37]

Future Directions in Recursive Medical AI

The future of the R-MoE framework lies in its ability to transition from 2D slice-based analysis to

fully volumetric 3D medical reasoning.[36] While current models excel at interpreting individual X-ray or MRI slices, they often fail to capture the spatial continuity required for identifying complex vascular structures or multi-planar orthopedic injuries.[36]

Research is also expanding the use of Reinforcement Learning from Human Feedback (RLHF) and Group Relative Policy Optimization (GRPO) to improve the "common sense" and ethical alignment of the reasoning agents.[25] By training models like Med-R1 on datasets of "Recursive Reasoning Paths," the AI community is building a shared foundation for more transparent and trustworthy clinical decision support systems.[7]

Ultimately, the R-MoE framework moves the industry closer to the goal of "Generalist Medical AI"—systems that can not only read images better than people but can also communicate context, manage clinical workflows, and maintain accountability alongside their human counterparts.[26] The shift toward modularity, recursion, and stochastic gating ensures that medical AI remains a safety-first technology, prioritizing patient outcomes over pure computational throughput.[15]

In conclusion, the Recursive Multi-Agent Mixture-of-Experts (R-MoE) for Autonomous Clinical Diagnostics represents a landmark in the evolution of medical technology. By decoupling the diagnostic process into specialized perception, reasoning, and reporting experts, and linking them through a recursive safety protocol, this system addresses the dual challenges of diagnostic hallucinations and visual ambiguity.[5] The empirical success of models like Qwen2-VL and DeepSeek-R1, combined with the structured clinical precision of Llama-3-Medius, provides a robust pathway toward a new era of highly accurate, interpretable, and safe autonomous radiology.[6] The ongoing refinement of recursive reasoning paths and the integration of human-in-the-loop workflows will ensure that these systems continue to enhance, rather than replace, the critical judgment of medical professionals in the decades to come.[30]

Works cited

1. Mixture of Experts in Large Language Models †: Corresponding author: Junhao Song (junhao.song23@imperial.ac.uk) - arXiv, accessed March 1, 2026, https://arxiv.org/html/2507.11181v1
2. Mixture of Experts in Large Language Models - ResearchGate, accessed March 1, 2026, https://www.researchgate.net/publication/393724282_Mixture_of_Experts_in_Large_Language_Models
3. Vision-Language Model Agent - Emergent Mind, accessed March 1, 2026, https://www.emergentmind.com/topics/vision-language-model-agent-5ca71059-e1f8-4036-9ab2-31569d49984c
4. Daily Papers - Hugging Face, accessed March 1, 2026, https://huggingface.co/papers?q=QWEN%20VL
5. Medical Reasoning in LLMs: An In-Depth Analysis of DeepSeek R1 - arXiv.org,

accessed March 1, 2026, https://arxiv.org/pdf/2504.00016
6. Knowledge-Guided Large Language Model for Automatic Pediatric Dental Record Understanding and Safe Antibiotic Recommendation - ResearchGate, accessed March 1, 2026, https://www.researchgate.net/publication/398560061_Knowledge-Guided_Large_Language_Model_for_Automatic_Pediatric_Dental_Record_Understanding_and_Safe_Antibiotic_Recommendation
7. MedVLM-R1: RL-Enhanced Medical VQA - Emergent Mind, accessed March 1, 2026, https://www.emergentmind.com/topics/medvlm-r1
8. Qwen2 VL 72B · Models - Dataloop, accessed March 1, 2026, https://dataloop.ai/library/model/qwen_qwen2-vl-72b/
9. Binomial Classification of Pediatric Elbow Fractures Using a Deep Learning Multiview Approach Emulating Radiologist Decision Making | Request PDF - ResearchGate, accessed March 1, 2026, https://www.researchgate.net/publication/330748156_Binomial_Classification_of_Pediatric_Elbow_Fractures_Using_a_Deep_Learning_Multiview_Approach_Emulating_Radiologist_Decision_Making
10. PathVLM-Eval: Evaluation of open vision language models in histopathology - PMC, accessed March 1, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12276438/
11. NOVA: A Benchmark for Rare Anomaly Localization and Clinical Reasoning in Brain MRI - OpenReview, accessed March 1, 2026, https://openreview.net/pdf/c906b235a22f7d5d108ff6d2aa15334ed7bc3792.pdf
12. Medical reasoning in LLMs: an in-depth analysis of DeepSeek R1 - PMC, accessed March 1, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12213874/
13. Medical Reasoning in LLMs: An In-Depth Analysis of DeepSeek R1 - ResearchGate, accessed March 1, 2026, https://www.researchgate.net/publication/390405366_Medical_Reasoning_in_LLMs_An_In-Depth_Analysis_of_DeepSeek_R1
14. Medical Reasoning in LLMs: An In-Depth Analysis of DeepSeek R1 - arXiv.org, accessed March 1, 2026, https://arxiv.org/html/2504.00016v1
15. Uncover This Tech Term: Uncertainty Quantification for Deep Learning - PMC, accessed March 1, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC10973738/
16. Uncertainty-Driven Expert Control: Enhancing the Reliability of Medical Vision-Language Models, accessed March 1, 2026, https://openaccess.thecvf.com/content/ICCV2025/papers/Liang_Uncertainty-Driven_Expert_Control_Enhancing_the_Reliability_of_Medical_Vision-Language_Models_ICCV_2025_paper.pdf
17. The challenge of uncertainty quantification of large language models in medicine - arXiv, accessed March 1, 2026, https://arxiv.org/html/2504.05278v1
18. DR-CoT: dynamic recursive chain of thought with meta reasoning for parameter efficient models - PMC, accessed March 1, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12501286/
19. Daily Papers - Hugging Face, accessed March 1, 2026, https://huggingface.co/papers?q=recursive%20model

20. Artificial Intelligence in Medical Imaging - MDPI, accessed March 1, 2026, https://mdpi-res.com/bookfiles/book/8264/Artificial_Intelligence_in_Medical_Imaging_The_Beginning_of_a_New_Era.pdf?v=1725152678
21. Uncertainty Estimation in Large Vision Language Models for Automated Radiology Report Generation - GitHub, accessed March 1, 2026, https://raw.githubusercontent.com/mlresearch/v259/main/assets/xu25a/xu25a.pdf
22. Evaluating Human–LLM Alignment Requires Transparent and Adaptable Statistical Guarantees - OpenReview, accessed March 1, 2026, https://openreview.net/pdf?id=iOSHFKHQNP
23. Measuring the relative contributions of rule-based and exemplar-based processes in judgment: Validation of a simple model | Judgment and Decision Making - Cambridge University Press & Assessment, accessed March 1, 2026, https://www.cambridge.org/core/journals/judgment-and-decision-making/article/measuring-the-relative-contributions-of-rulebased-and-exemplarbased-processes-in-judgment-validation-of-a-simple-model/A716B1E97BA9F44D92D5FA3897521BA7
24. A Lightweight Framework Integrating MoE and SAM with a Sparse Attention Mechanism for Plant Disease Segmentation in Resource-Constrained Environments - MDPI, accessed March 1, 2026, https://www.mdpi.com/2223-7747/14/17/2634
25. Critique of impure reason: Unveiling the reasoning behaviour of medical large language models | eLife, accessed March 1, 2026, https://elifesciences.org/articles/106187
26. What is Human in the Loop (HITL)? Complete Guide 2026 - Articsledge, accessed March 1, 2026, https://www.articsledge.com/post/human-in-the-loop-hitl
27. Knowledge-Guided Large Language Model for Automatic Pediatric Dental Record Understanding and Safe Antibiotic Recommendation - arXiv, accessed March 1, 2026, https://www.arxiv.org/pdf/2512.09127
28. Large Language Models: A Comprehensive Guide for ... - PMC - NIH, accessed March 1, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC11473987/
29. Stanford Develops Real-World Benchmarks for Healthcare AI Agents, accessed March 1, 2026, https://www.makebot.ai/blog-en/healthcare-ai-agent-benchmark
30. 1: Beyond the Hype: How Large Language Models Are Empowering, Not Replacing, Medical Practitioners - ML-Neuro - BayernCollab, accessed March 1, 2026, https://collab.dvb.bayern/spaces/TUMmlneuro/pages/1242999257/%E2%AD%901+Beyond+the+Hype+How+Large+Language+Models+Are+Empowering+Not+Replacing+Medical+Practitioners
31. Guide to Human-in-the-Loop Machine Learning - Appen, accessed March 1, 2026, https://www.appen.com/blog/human-in-the-loop
32. View Session Recordings: Virtual AI Summit - Beyond the Hype - Concurrency, accessed March 1, 2026, https://concurrency.com/events/virtual-ai-summit-beyond-the-hype/
33. Mixture of Experts in Large Language Models †: Corresponding author - arXiv, accessed March 1, 2026, https://arxiv.org/html/2507.11181v2

34. Track: San Diego Poster Session 5 - NeurIPS, accessed March 1, 2026, https://neurips.cc/virtual/2025/loc/san-diego/session/128335
35. Med-R1: Reinforcement Learning for Generalizable Medical Reasoning in Vision-Language Models - arXiv.org, accessed March 1, 2026, https://arxiv.org/html/2503.13939v1
36. MedVL-SAM2: A unified 3D medical vision–language model for multimodal reasoning and prompt-driven segmentation - arXiv, accessed March 1, 2026, https://arxiv.org/html/2601.09879v1
37. In-context Learning for Zero-shot Medical Report Generation - OpenReview, accessed March 1, 2026, https://openreview.net/pdf/a9638ad85078f8e3ce12211d18a4382234fa219e.pdf
38. Daily Papers - Hugging Face, accessed March 1, 2026, https://huggingface.co/papers?q=vision-language%20models%20(VLMs)
39. Implementing generative artificial intelligence in precision oncology: safety, governance, and significance - PMC, accessed March 1, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12896320/
40. [Literature Review] MedVLM-R1: Incentivizing Medical Reasoning Capability of Vision-Language Models (VLMs) via Reinforcement Learning - Moonlight, accessed March 1, 2026, https://www.themoonlight.io/en/review/medvlm-r1-incentivizing-medical-reasoning-capability-of-vision-language-models-vlms-via-reinforcement-learning
41. [Literature Review] Med-R1: Reinforcement Learning for Generalizable Medical Reasoning in Vision-Language Models - Moonlight, accessed March 1, 2026, https://www.themoonlight.io/en/review/med-r1-reinforcement-learning-for-generalizable-medical-reasoning-in-vision-language-models
42. Reasoning Beyond Limits: Advances and Open Problems for LLMs - arXiv.org, accessed March 1, 2026, https://arxiv.org/html/2503.22732v2
43. I am a high level engineer. AI can already mostly replace my job. - Reddit, accessed March 1, 2026, https://www.reddit.com/r/accelerate/comments/1ljt1bn/i_am_a_high_level_engineer_ai_can_already_mostly/
44. Mixture of Experts in Large Language Models - arXiv.org, accessed March 1, 2026, https://arxiv.org/pdf/2507.11181?
45. Daily Papers - Hugging Face, accessed March 1, 2026, https://huggingface.co/papers?q=Medical%20benchmarks